



RNA-seq to study HIV Infection in cells

Jason Laird

Bioinformatics Scientist

Feb 2022

Research Technology Team



Delilah Maloney
High Performance Computing Specialist



Kyle Monahan
Senior Data Science Specialist



Shawn Doughty
Manager, Research Computing



Jason Laird
Bioinformatics Scientist



Chris Barnett
Senior Geospatial Analyst



Tom Phimmasen
Senior Data Consultant



Patrick Florance
Director, Academic Data Services



Jake Perl
Digital Humanities NLP Specialist



Carolyn Talmadge
Senior GIS Specialist



Uku-Kaspar Uustalu
Data Science Specialist

- ✓ Consultation on Projects and Grants
- ✓ High Performance Compute Cluster
- ✓ Workshops

<https://it.tufts.edu/research-technology>

Outline

Bulk and single cell
RNA sequencing

Intro to Galaxy Platform for
Bioinformatics (Tufts network or
VPN required)


<https://galaxy.cluster.tufts.edu/>

Work through RNAseq
example together on Galaxy

[https://rbatorsky.github.io/in
tro-to-rnaseq-with-galaxy/](https://rbatorsky.github.io/intro-to-rnaseq-with-galaxy/)

2 days!


DNA and RNA in a cell



Two common analyses

DNA Sequencing

- Fixed number of copies of a gene per cell
- Analysis goal:
Variant calling and interpretation




RNA Sequencing

- Number of copies of a gene transcript per cell depends on gene expression
- Analysis goal:
 - Bulk : Differential expression
 - Single cell : Quantify different cell populations

Today we will cover RNA sequencing

DNA Sequencing

- Fixed number of copies of a gene per cell
- Analysis goal: Variant calling and interpretation



RNA Sequencing

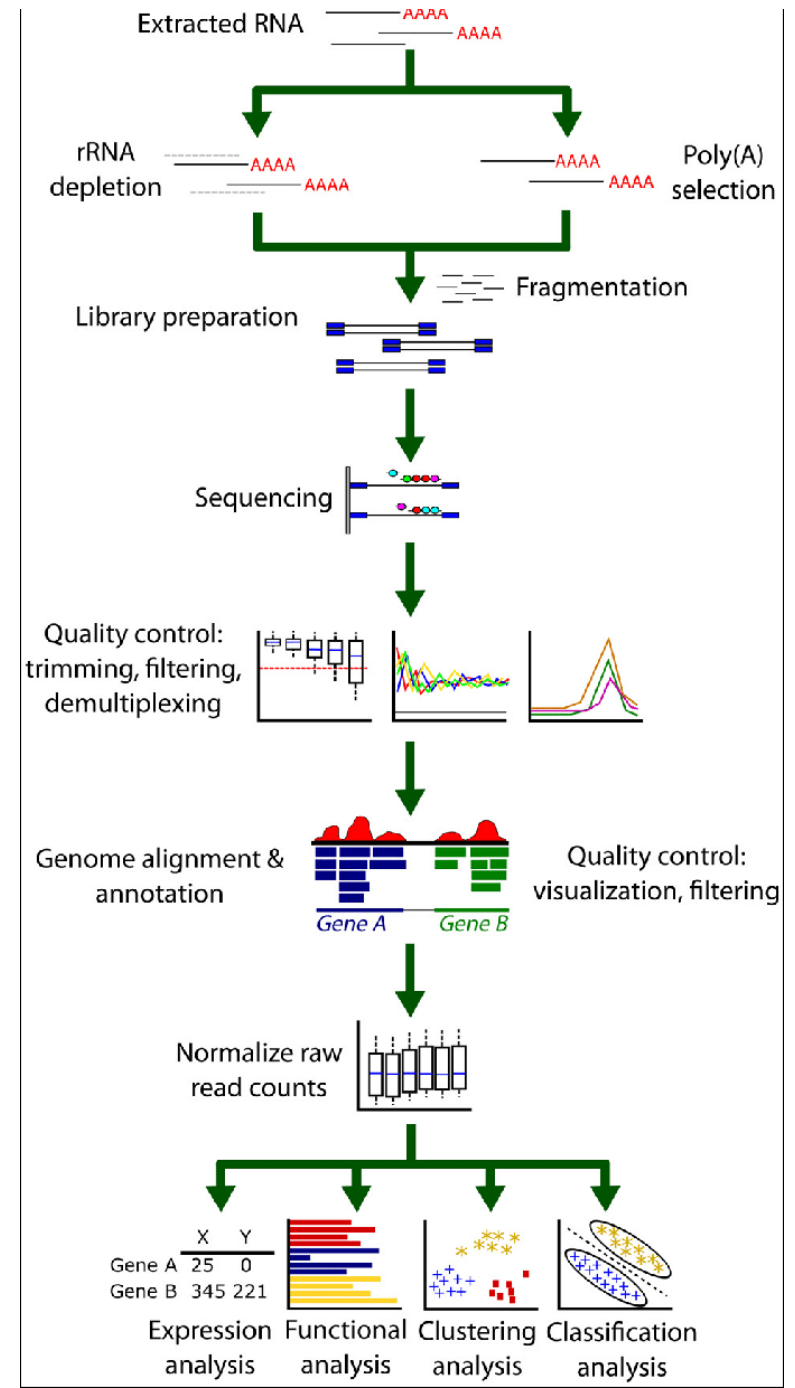
- Number of copies of a gene transcript per cell depends on gene expression
- Analysis goal:
 - Bulk : Differential expression
 - Single cell : Quantify different cell populations

“Bulk” RNA seq workflow


Library prep and sequencing

Bioinformatics

Good resource: [Griffiths et al Plos Comp Bio 2015](#)

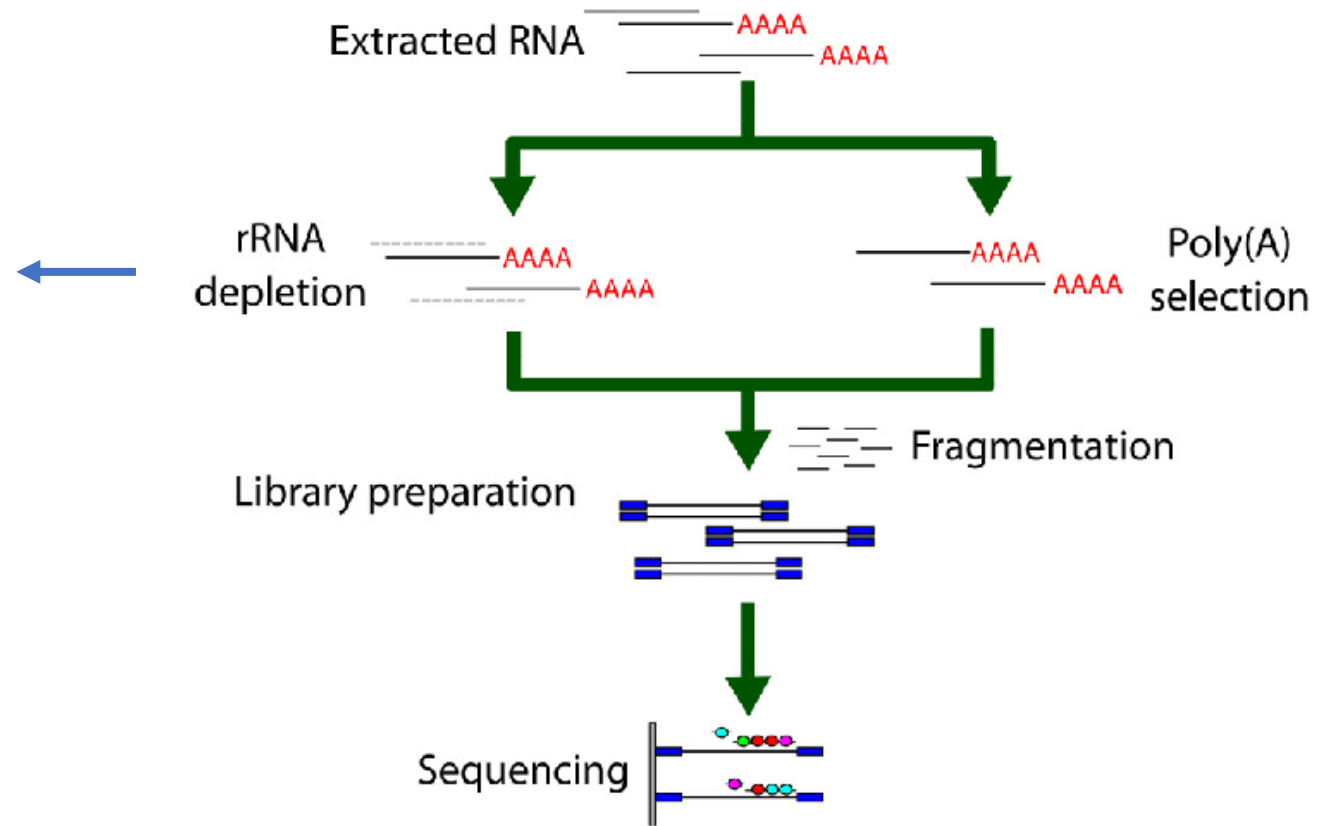


RNA seq library prep and sequencing



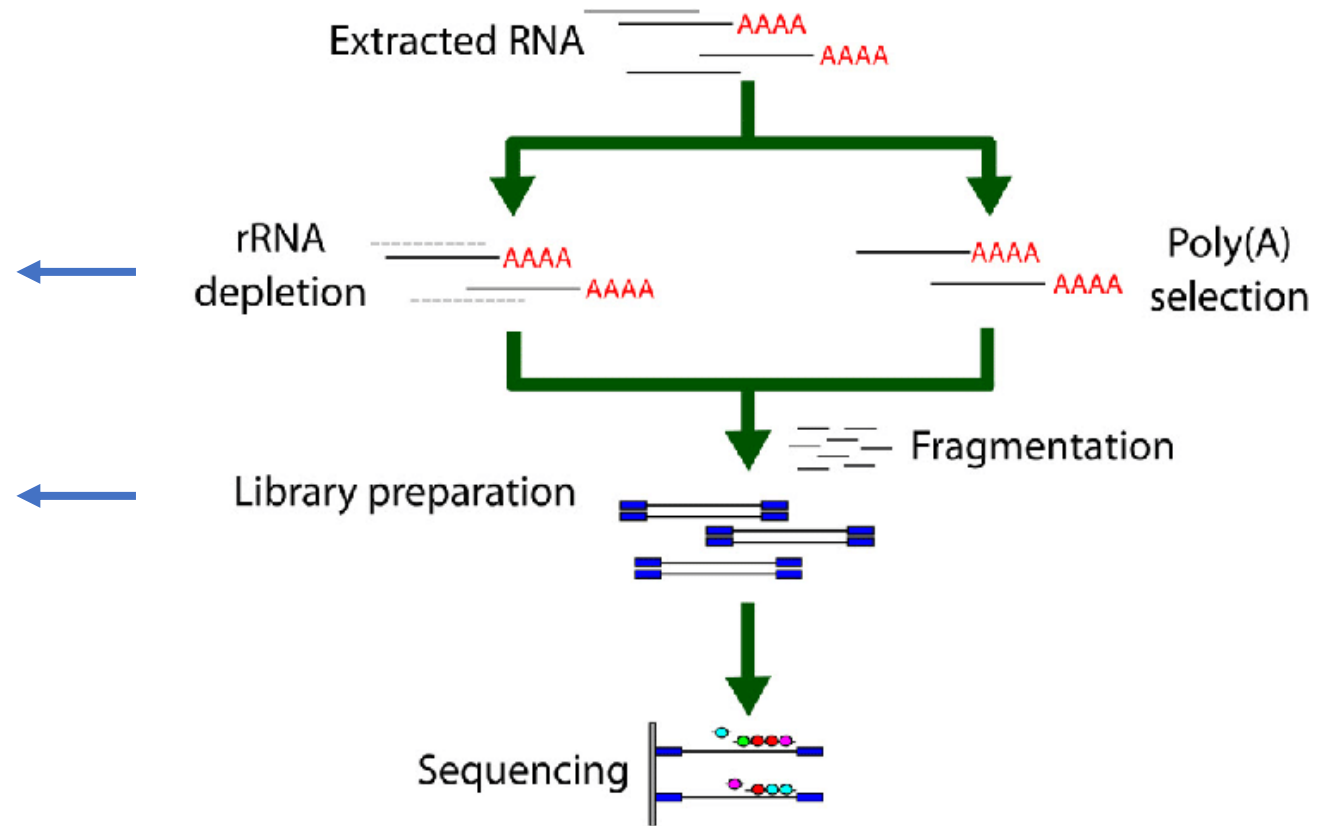
RNA seq library prep and sequencing

- Enrichment for mRNA, two options
- In humans, ~95%–98% of all RNA molecules are rRNAs



RNA seq library prep and sequencing

- Enrichment for mRNA, two options
- In humans, ~95%–98% of all RNA molecules are rRNAs
- Random priming and reverse transcription
- Double stranded cDNA synthesis
- Sequencing adapter ligation



Resources:

[Illumina Sequencing by Synthesis](#)

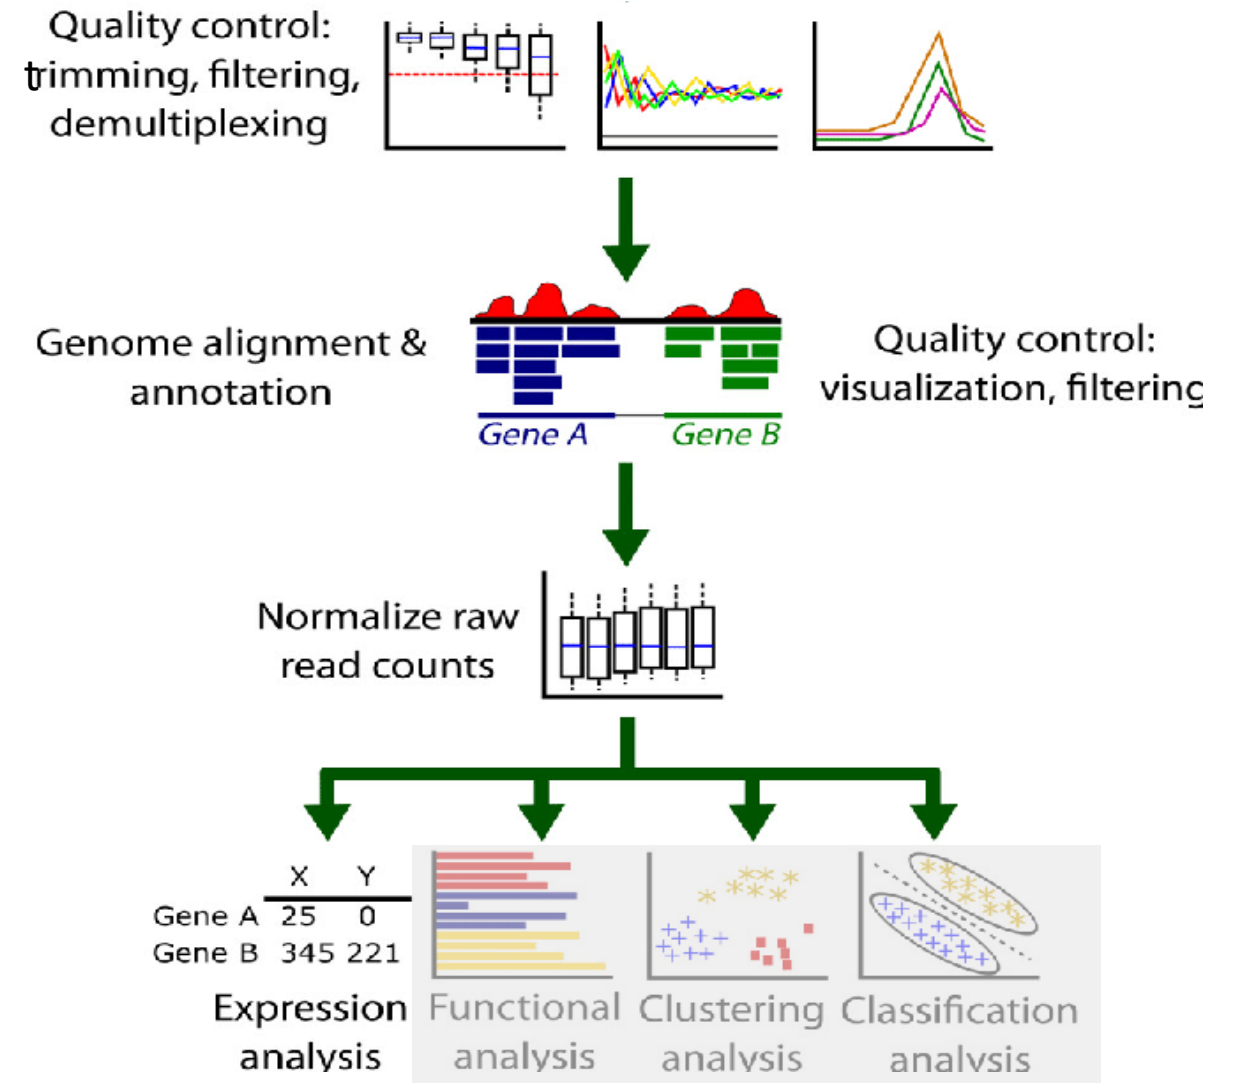
[Griffiths et al Plos Comp Bio 2015](#)

RNA seq bioinformatics

Goal of Differential Expression

“How can we detect genes for which the counts of reads change between conditions **more systematically** than as expected by chance”

Oshlack et al. 2010. From RNA-seq reads to differential expression results. Genome Biology 2010, 11:220




Our dataset


Next-Generation Sequencing Reveals HIV-1-Mediated Suppression of T Cell Activation and RNA Processing and Regulation of Noncoding RNA Expression in a CD4⁺ T Cell Line

Stewart T. Chang, Pavel Sova, Xinxia Peng, Jeffrey Weiss, G. Lynn Law, Robert E. Palermo, Michael G. Katze

Mock Infected
CD4⁺ T Cells




HIV Infected
CD4⁺ T Cells



12 hour



24 hour

HIV lifecycle





HIV lifecycle

HIV infection in a human host



The study question

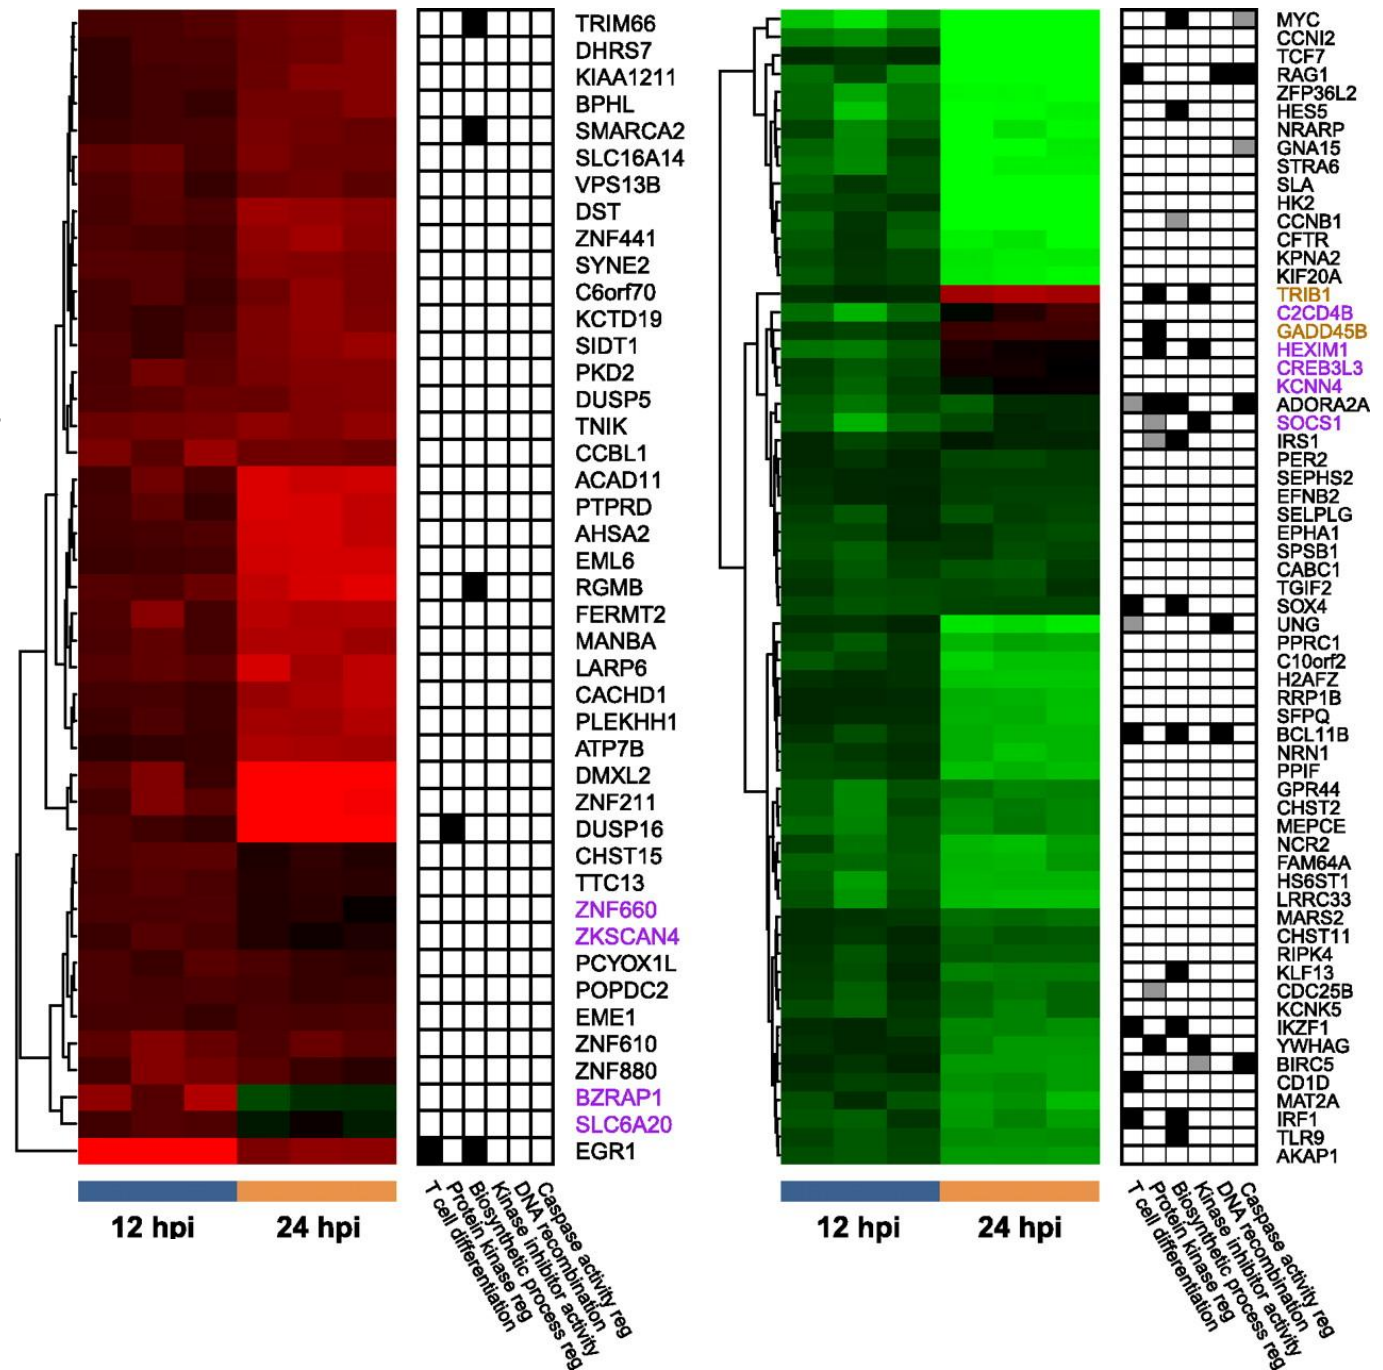

What changes take place in the first 12-24 hours of HIV infection in terms of gene expression of host cell and viral replication levels?



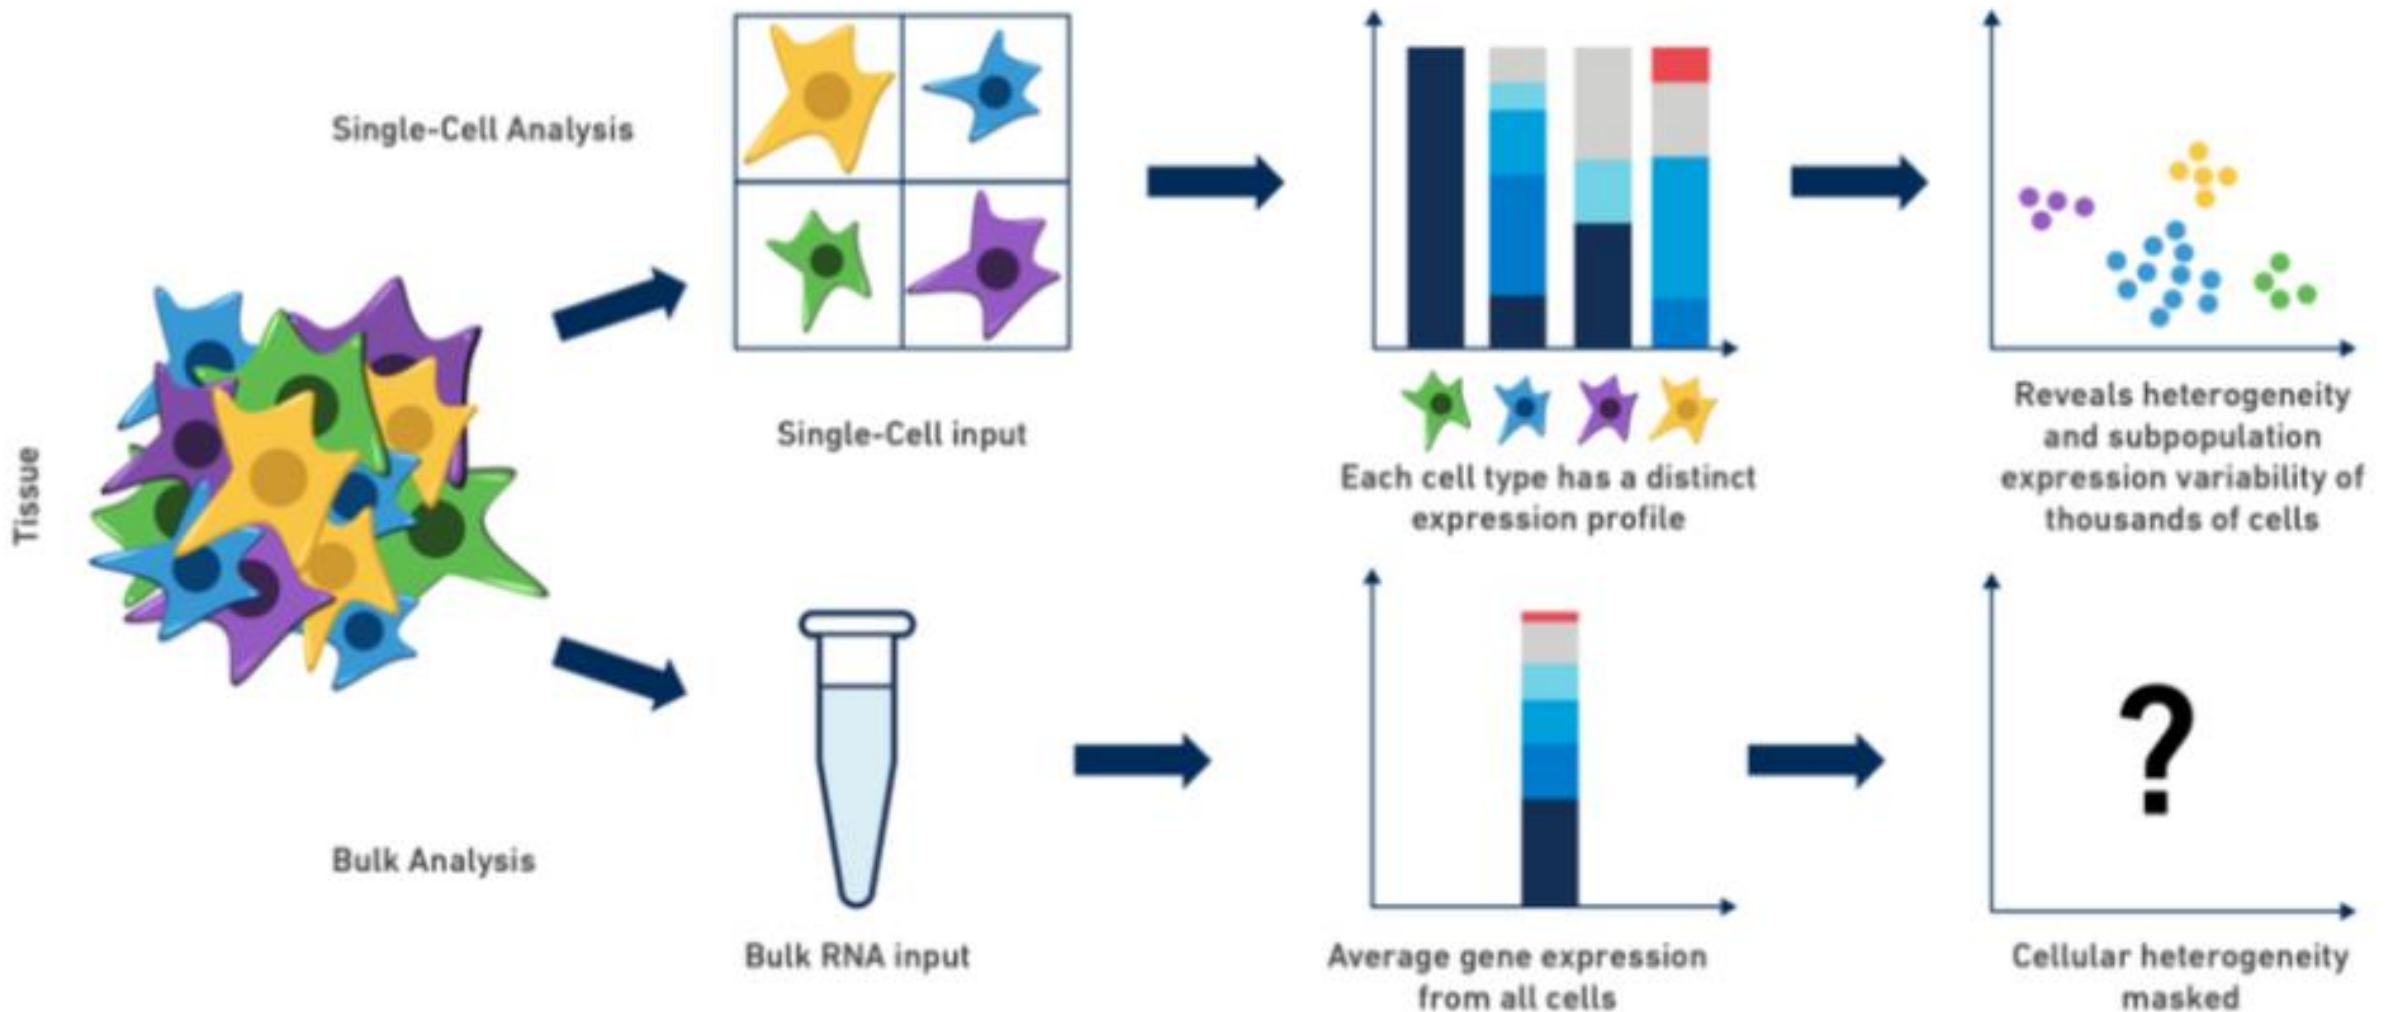
Study findings

Using RNAseq, authors demonstrate:


- 20% of reads mapped to HIV at 12 hr, 40% at 24hr
- Downregulation of T cell differentiation genes at 12hr
- ‘Large-scale disruptions to host transcription’ at 24hr




Bulk vs Single Cell RNA Sequencing




scRNA cell subsets in PBMC



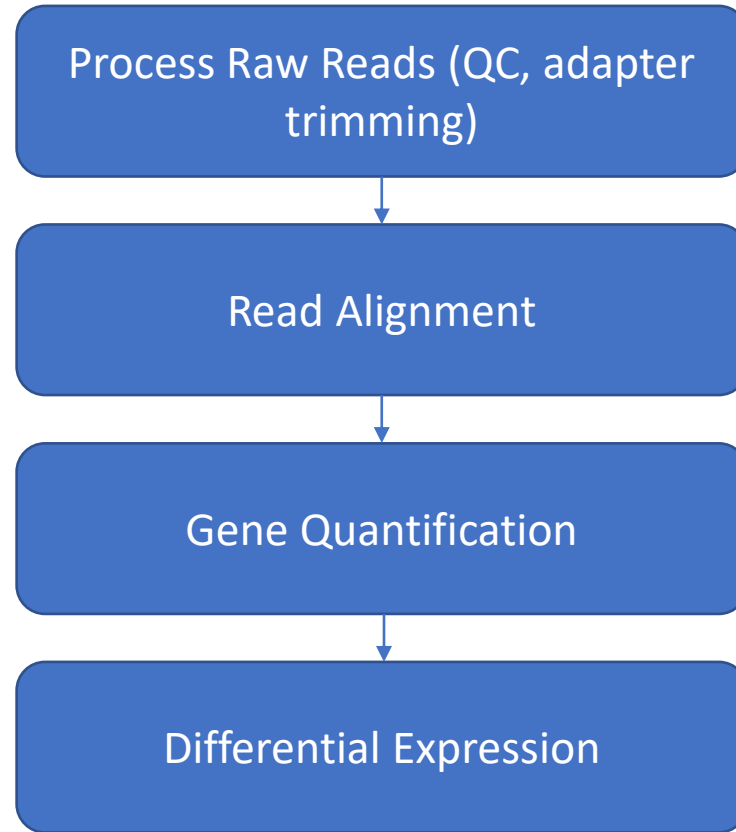
10x single cell technology



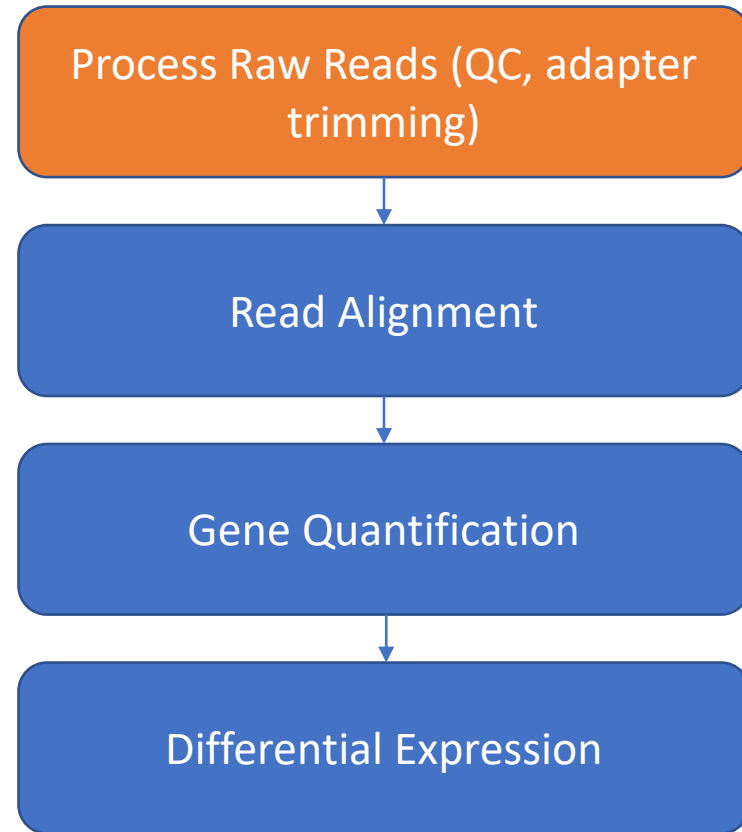
Bulk RNAseq for Differential Expression is OK!



Our (bulk) RNAseq Workflow



Quality control on Raw Reads



Raw reads in Fastq format

```
@SRR098401.109756285  
GACTCACGTAAC TTAACTCTAACAGAAATATACTA...  
+  
CAEFGDG?BCGGGEEDGGHGHGDFHEIEGGDDDD...
```

1. Sequence identifier
2. Sequence
3. + (optionally lists the sequence identifier again)
4. Quality string

Base Quality Scores

The symbols we see in the read quality string are an encoding of the quality score:


```
Quality encoding: !"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHI
                  |           |           |           |
Quality score: 0.....10.....20.....30.....40
```

A quality score is a prediction of the probability of an error in base calling:

Quality Score	Probability of Incorrect Base Call	Inferred Base Call Accuracy
10 (Q10)	1 in 10	90%
20 (Q20)	1 in 100	99%
30 (Q30)	1 in 1000	99.9%

Base Quality Scores

The symbols we see in the read quality string are an encoding of the quality score:



A quality score is a prediction of the probability of an error in base calling:

Quality Score	Probability of Incorrect Base Call	Inferred Base Call Accuracy
10 (Q10)	1 in 10	90%
20 (Q20)	1 in 100	99%
30 (Q30)	1 in 1000	99.9%

Back to our read:

```
@SRR098401.109756285
GACTCACGTAAC TTAACTCTAACAGAAATATACTA...
+
CAEFGDG?BCGGGEEDGGHGHGDFHEIEGGDDDD...
```

↑ C → Q = 34 → Probability < 1/1000 of an error

Raw read quality control

Fastq File

```
@SRR497699.30343179.1 HWI-EAS39X_10175_FC61MK0_4_117_4812_10346 length=75
CAGATGGCCGCAGAGGAAGCCATGAAGGCCCTGCATGGGGAGATCGGAAGAGCGTTTCAGCAGGAATGCCGAGAC
+
IIIIIGIIHFIIIIIBIIDII>IIDHIIHDIIIGIFIIEIGIBDDEFIG<EIEGEEG;<DB@A8CC7<><C@BBDDDB
@SRR497699.11626500.1 HWI-EAS39X_10175_FC61MK0_4_44_8384_16550 length=75
CGTACTGAACGTACAACGCTGATGCCATCCGCATATTTAAATTCGGCAGCGTTAATTAACCTCCCTGACCTCGGCG
+
HHHHHHHHHHFHHHGHHHHHB@HHHHHHHHFHHHHEHHHHHHHHHHHGEHDHHEHHHHBHHHGHHHHHHHG
@SRR497699.29057557.1 HWI-EAS39X_10175_FC61MK0_4_112_12508_19308 length=75
CCGAGGCTTAGCTTTCATTATCACTGTCTCCAGGGTGTGCTTGTCAAAGAGATAAGATCGGAAGAGCGGTTTCAG
+
GGGBGGGDGBHHDHHEGGGHHHHHGHGHGHHHHHGBGGDGGEGDHHHHHHHHHHH@BHHGGHGHHHHHEEGHH
@SRR497699.1331889.1 HWI-EAS39X_10175_FC61MK0_4_5_4738_15920 length=75
CTTACTTTGTAGCCTTCATCAGGGTTTGCTGAAGATGGCGGTATATAGGCTGAGCAAGAGGTGGTGAGGTTGATC
+
HHHHHHHHHHGGGGGHHGHGEBEEGGEDGGGGGHHHHHGGEGBDGGGDDGBGGC<EADBEBE<GGGGBEEDGD
```

...


FastQC Tool



Metrics

- Sequence Quality
- GC content
- Per base sequence content
- Adapters in Sequence

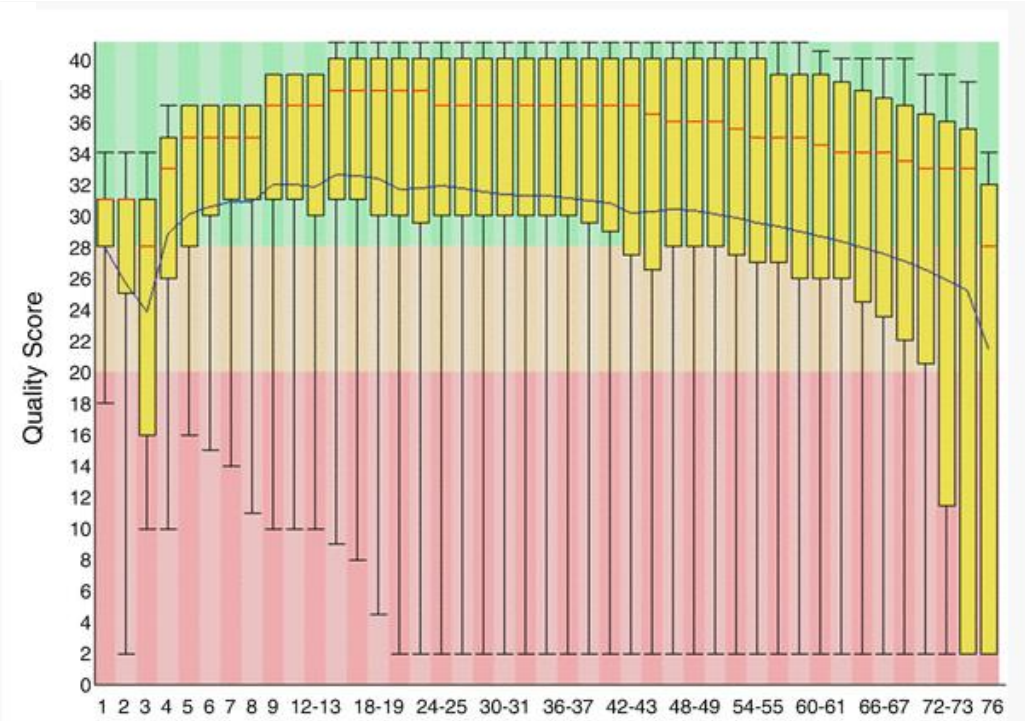
FastQC: Sequence Quality Histogram



Position in read (bp)

GOOD


High quality over the length of the read



Position in read (bp)



BAD

Read quality drops at the beginning and end




FastQC: Per sequence GC content

✓ Per sequence GC content




FastQC: Per sequence GC content

✓ Per sequence GC content




GOOD: follows normal distribution (sum of deviations is < 15% of reads)

✗ Per sequence GC content




BAD: can indicate contamination with adapter dimers, or another species

FastQC: Per Base Sequence Content




- Proportion of each position for which each DNA base has been called
- RNAseq data tends to show a positional sequence bias in the first ~12 bases
- The "random" priming step during library construction is not truly random and certain hexamers are more prevalent than others

FastQC: Per Base Sequence Content




EXPECTED for RNAseq



BAD:

Shows a strong positional bias throughout the reads, which in this case is due to the library having a certain sequence that is overrepresented


FastQC: Adapter content



FastQC will scan each read for the presence of known adapter sequences


The plot shows that the adapter content rises over the course of the read

Solution – Adapter trimming!



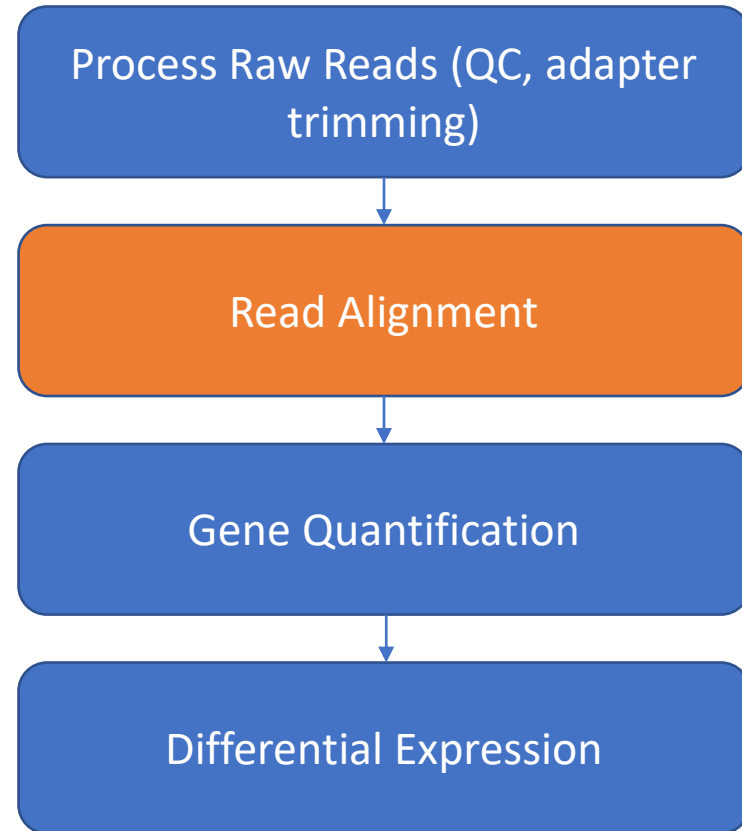
FastQC -> MultiQC

Should view all samples at once to notice abnormalities for our dataset.



We'll use a tool called "Trim Galore!" to trim adapters and remove low quality bases/reads.

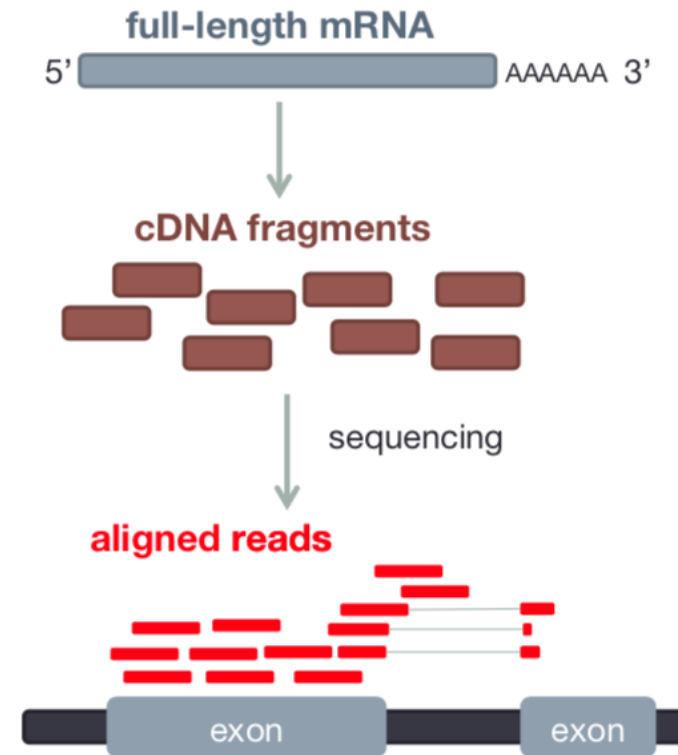
Workflow




Read Alignment

- RNAseq data originates from spliced mRNA (no introns)
- When aligning to the genome, our aligner must find a spliced alignment for reads
- We use a tool called STAR (Spliced Transcripts Alignment to a Reference) that has an exon-aware mapping algorithm.

Reference sequence



Sequence Alignment Map (SAM)



@HD VN:1.5 SO:coordinate										
@SQ SN:ref LN:45										
r001	99	ref	7	30	8M2I4M1D3M	=	37	39	TTAGATAAAGGATACTG	*
r002	0	ref	9	30	3S6M1P1I4M	*	0	0	AAAAGATAAGGATA	*
r003	0	ref	9	30	5S6M	*	0	0	GCCTAAGCTAA	* SA:Z:ref,29,-,6H5M,17,0;
r004	0	ref	16	30	6M14N5M	*	0	0	ATAGCTTCAGC	*
r003	2064	ref	29	17	6H5M	*	0	0	TAGGC	* SA:Z:ref,9,+,5S6M,30,1;
r001	147	ref	37	30	9M	=	7	-39	CAGCGGCAT	* NM:i:1

Header
section

Alignment
section

↑ CIGAR: summary of alignment, e.g. match, gap, insertion, deletion

↑ Mapping Quality


↑ Position

↑ Ref Sequence name

↑ Flag: indicates alignment information e.g. paired, aligned, etc
<https://broadinstitute.github.io/picard/explain-flags.html>

↑ Read ID

Sequence Alignment Map (SAM)



@HD VN:1.5 SO:coordinate										
@SQ SN:ref LN:45										
r001	99	ref	7	30	8M2I4M1D3M	=	37	39	TTAGATAAAGGATACTG	*
r002	0	ref	9	30	3S6M1P1I4M	*	0	0	AAAAGATAAGGATA	*
r003	0	ref	9	30	5S6M	*	0	0	GCCTAAGCTAA	* SA:Z:ref,29,-,6H5M,17,0;
r004	0	ref	16	30	6M14N5M	*	0	0	ATAGCTTCAGC	*
r003	2064	ref	29	17	6H5M	*	0	0	TAGGC	* SA:Z:ref,9,+,5S6M,30,1;
r001	147	ref	37	30	9M	=	7	-39	CAGCGGCAT	* NM:i:1

Header
section

Alignment
section

Paired end info


Sequence

Quality Score

Optional Fields

Genome Annotation Standards

- STAR can use an annotation file gives the location and structure of genes in order to improve alignment in known splice junctions
- Annotation is dynamic and there are at least three major sources of annotation
- The intersection among RefGene, UCSC, and Ensembl annotations shows high overlap. RefGene has the fewest unique genes, while more than 50% of genes in Ensembl are unique
- Be consistent with your choice of annotation source!

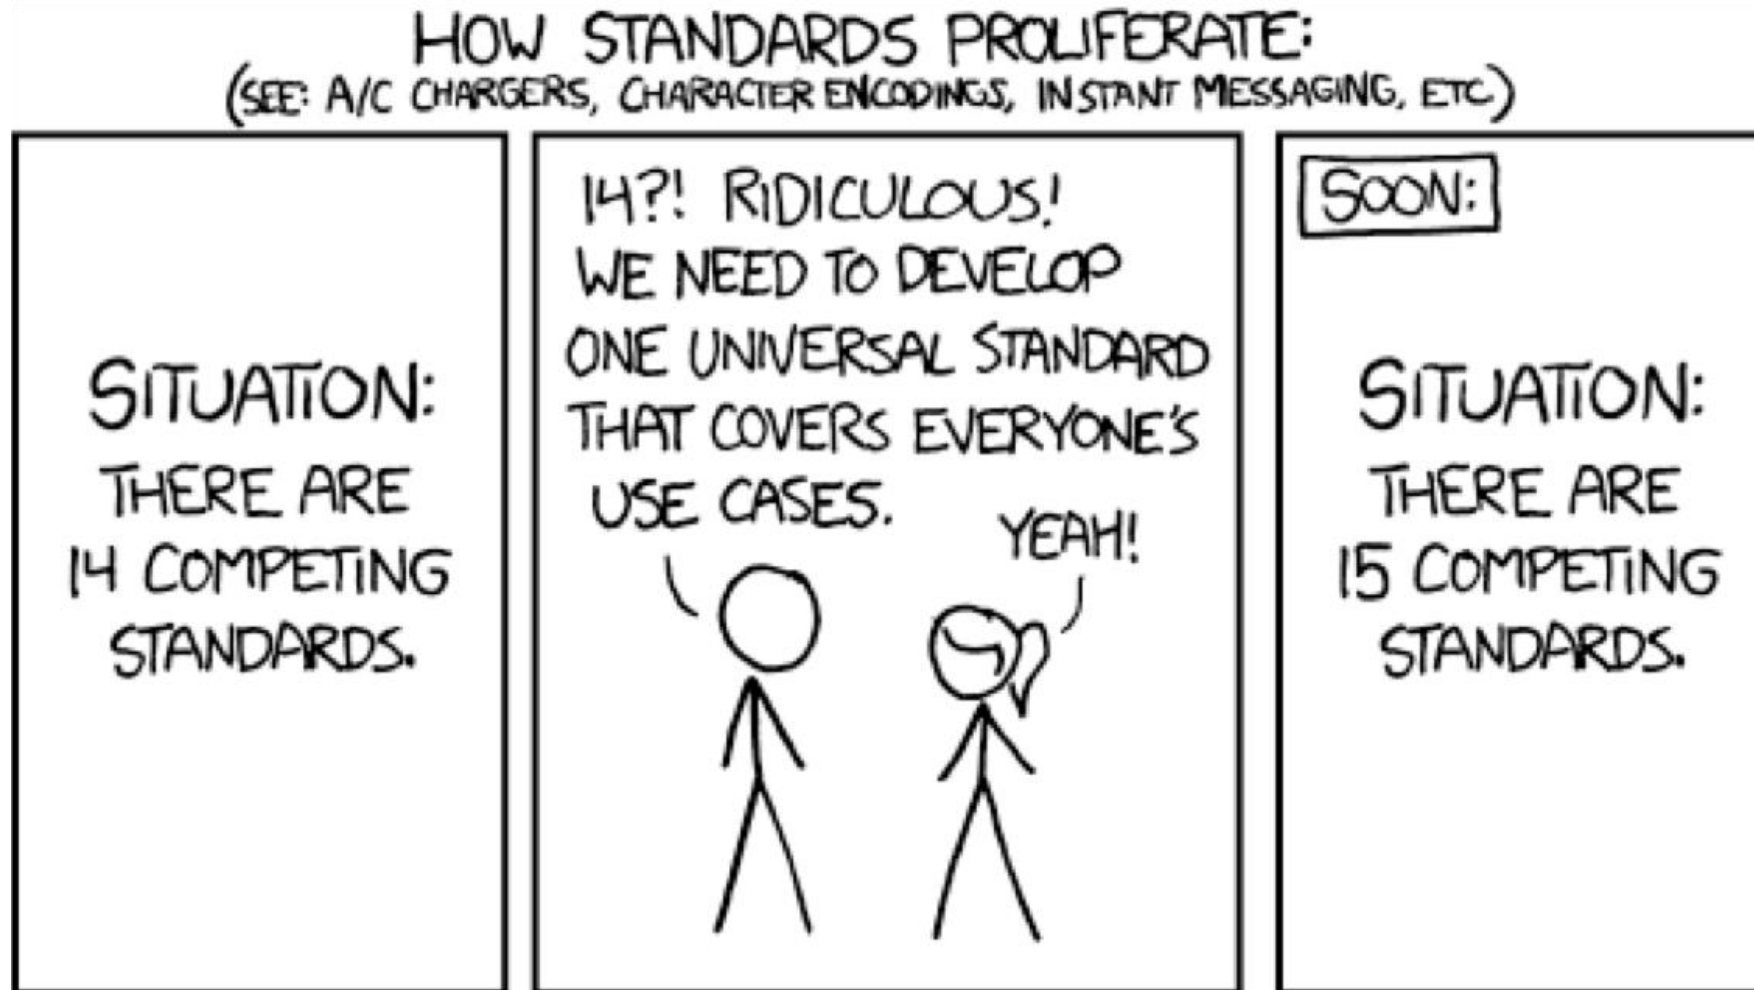


Gene Annotation Format (GTF)

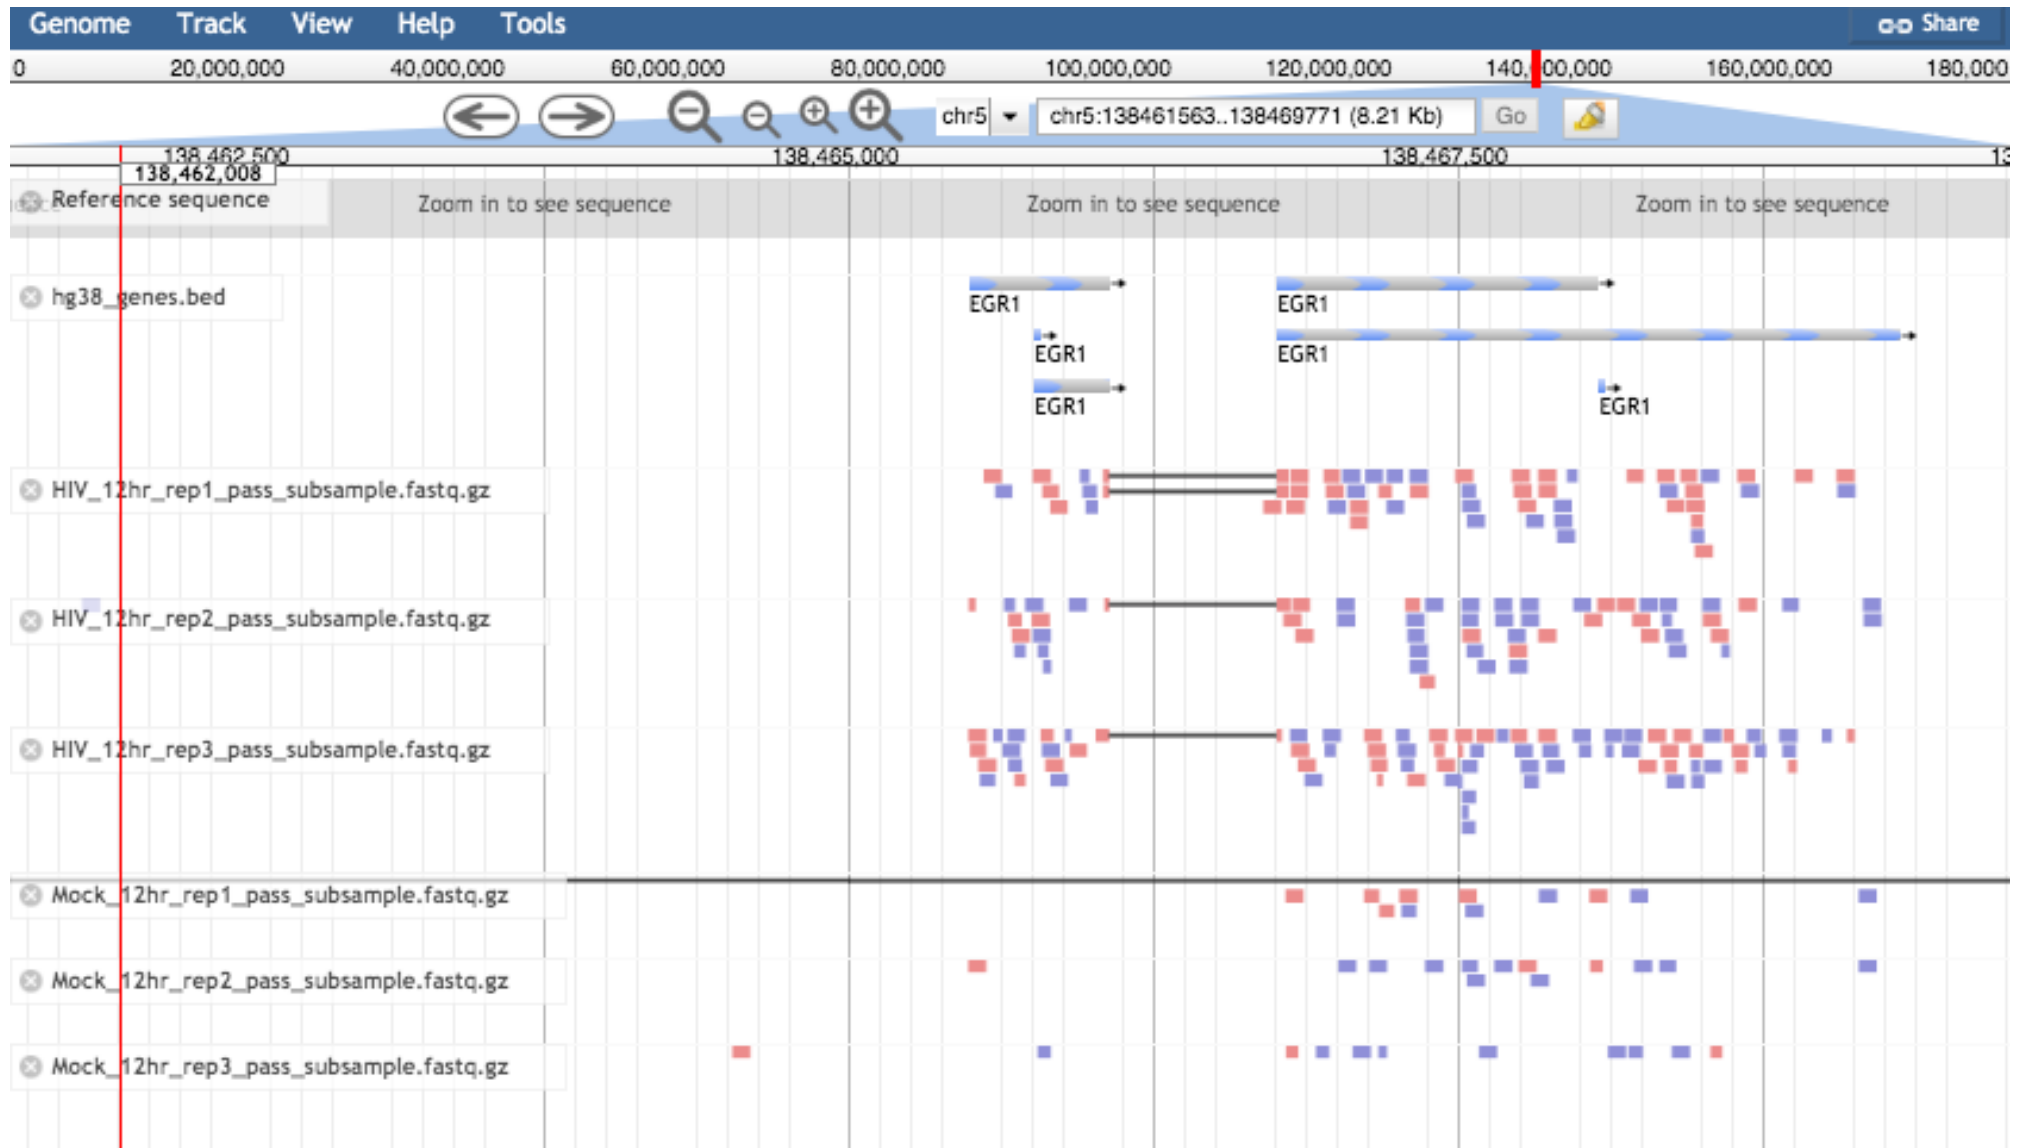
In order to count genes, we need to know where they are located in the reference sequence
STAR uses a Gene Transfer Format (GTF) file for gene annotation

Chrom	Source	Feature type	Start	Stop	Strand (Score) Frame			Attribute
chr5	hg38_refGene	exon	138465492	138466068	.	+	.	gene_id "EGR1";
chr5	hg38_refGene	CDS	138465762	138466068	.	+	0	gene_id "EGR1";
chr5	hg38_refGene	start_codon	138465762	138465764	.	+	.	gene_id "EGR1";
chr5	hg38_refGene	CDS	138466757	138468078	.	+	2	gene_id "EGR1";
chr5	hg38_refGene	exon	138466757	138469315	.	+	.	gene_id "EGR1";
chr5	hg38_refGene	stop_codon	138468079	138468081	.	+	.	gene_id "EGR1";

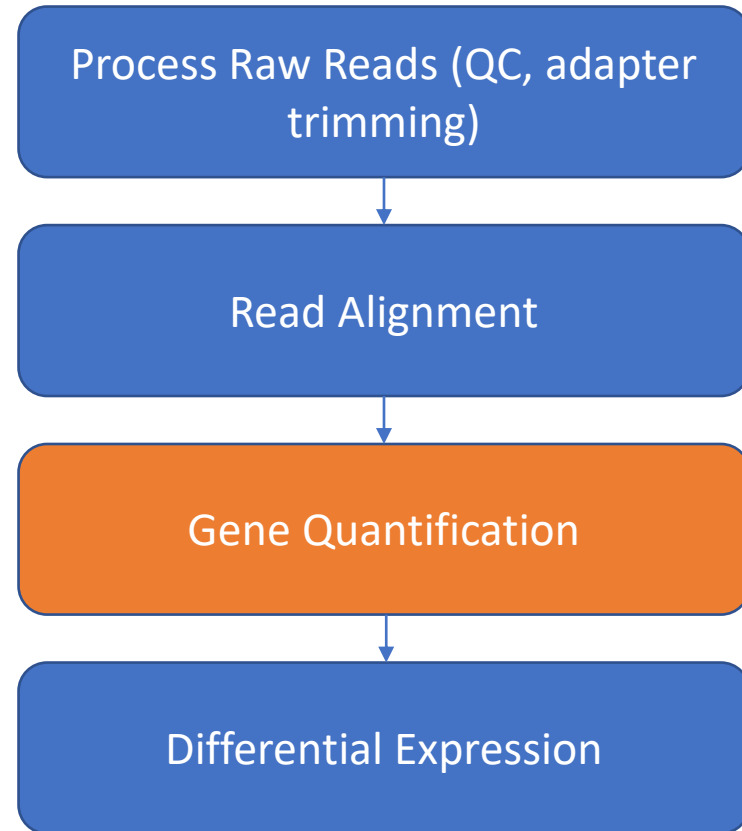
A note on standards




Visualizing reads with JBrowse



Workflow

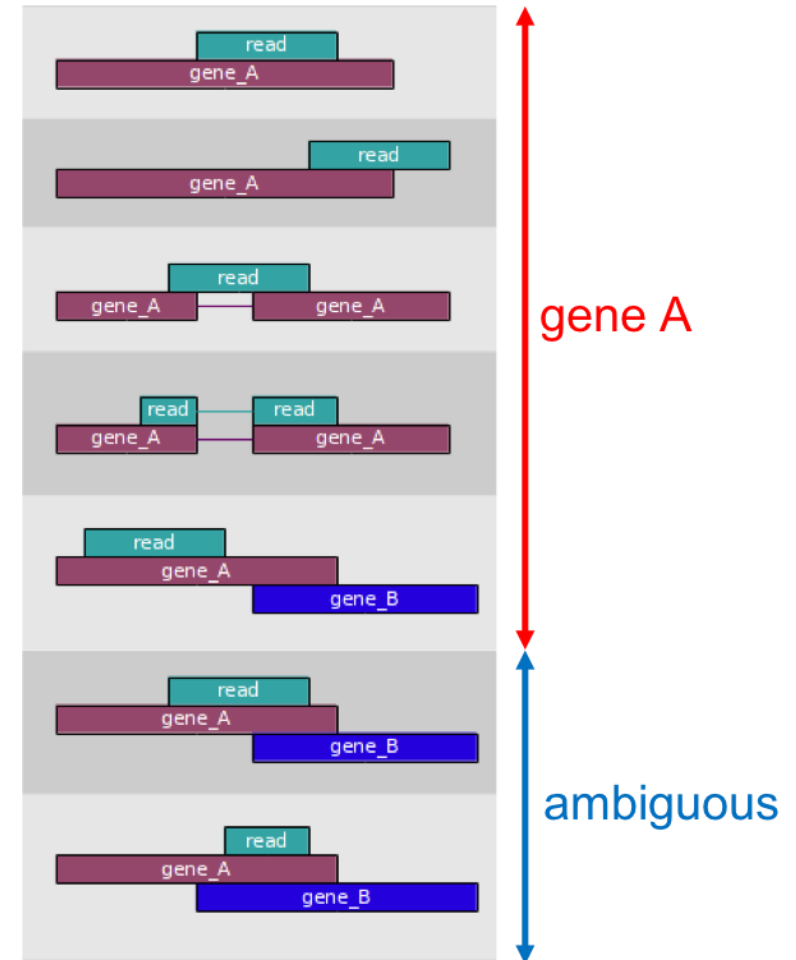


Counting reads for each gene



Counting reads: featurecounts

- The mapped coordinates of each read are compared with the features in the GTF file
- Reads that overlap with a gene by ≥ 1 bp are counted as belonging to that feature
- Ambiguous reads will be discarded

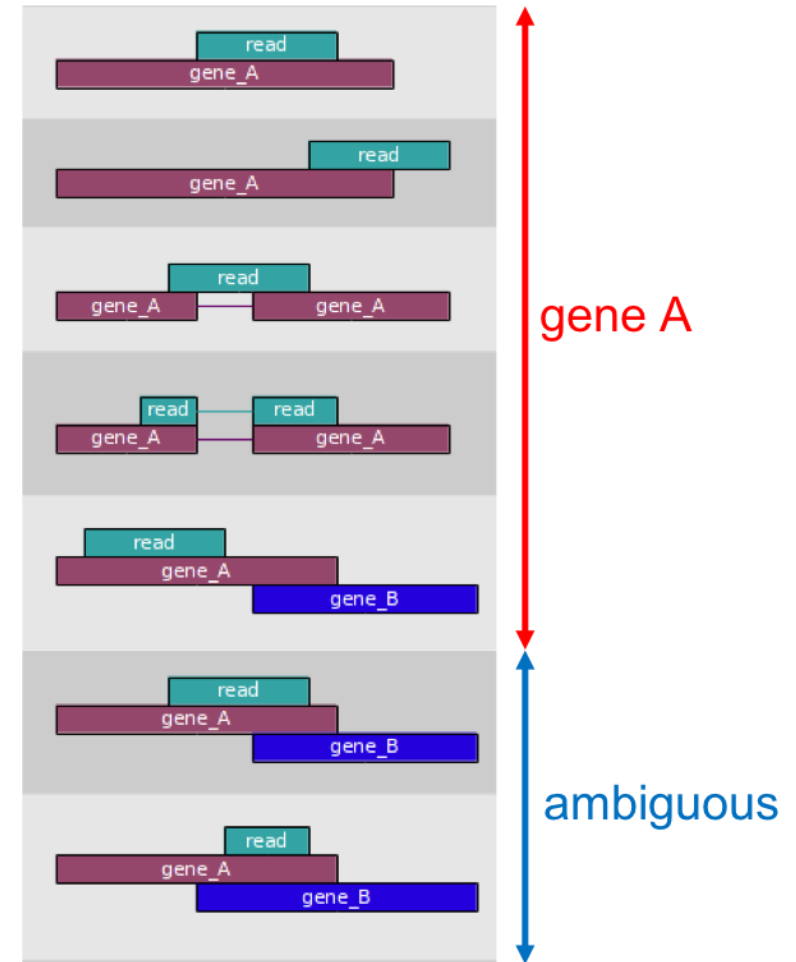


Counting reads: featurecounts

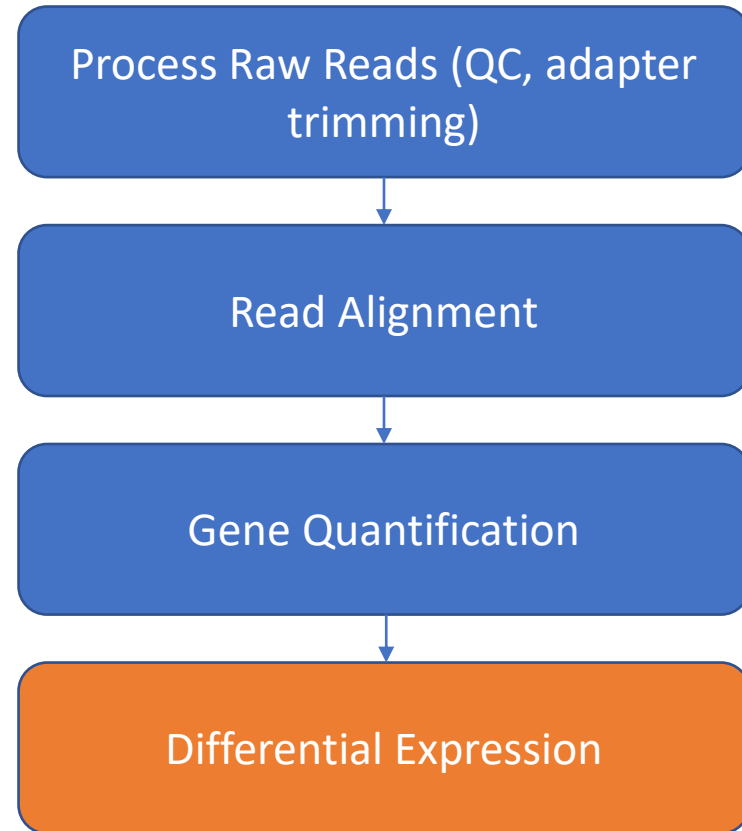
- The mapped coordinates of each read are compared with the features in the GTF file
- Reads that overlap with a gene by ≥ 1 bp are counted as belonging to that feature
- Ambiguous reads will be discarded

Result is a gene count matrix:

Gene	Sample 1	Sample 2	Sample 3	Sample 4
A	1000	1000	100	10
B	10	1	5	6
C	10	1	10	20




Workflow



Normalization

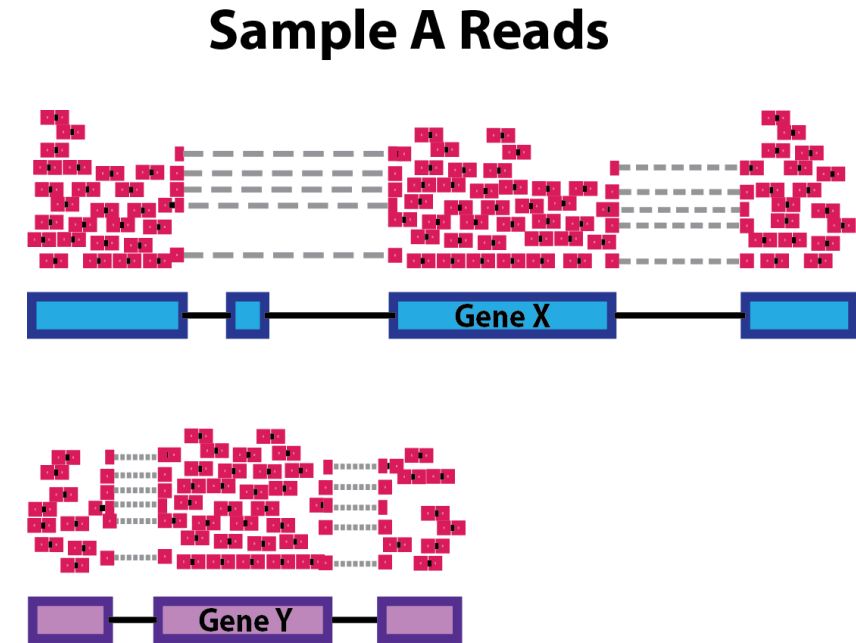
- Raw Count \neq Expression strength
- Normalization:
 - Eliminates factors that are not of interest for our experiment
 - Enables accurate comparison between samples or genes



Normalization

The number of reads mapped to a gene depends on

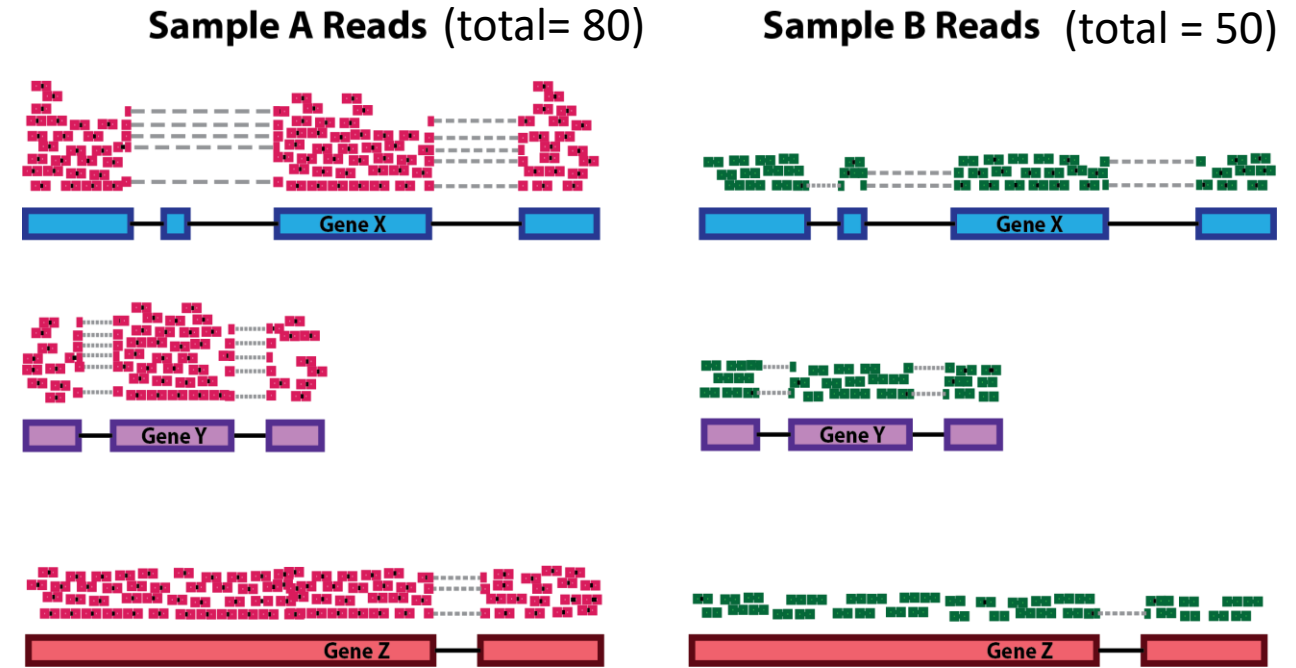
- **Gene Length**



Normalization

The number of reads mapped to a gene depends on

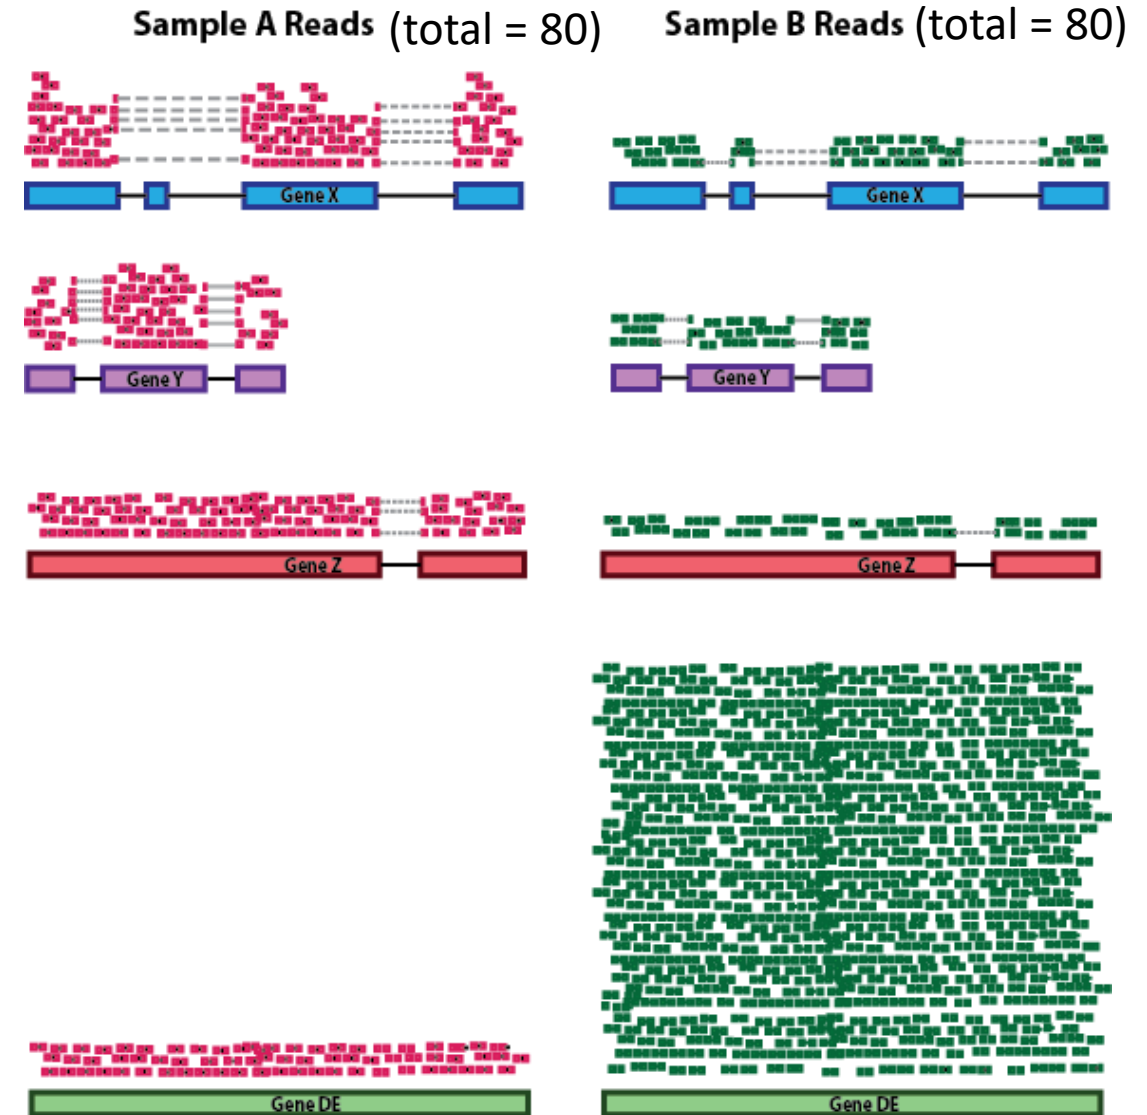
- Gene Length
- **Sequencing depth**



Normalization

The number of reads mapped to a gene depends on

- Gene Length
- Sequencing depth
- **The expression level of other genes in the sample (RNA Composition)**




Normalization

The number of reads mapped to a gene depends on

- Gene Length
- Sequencing depth
- The expression level of other genes in the sample (RNA Composition)

DESeq2 Median of Ratios




Normalization: DESeq2 Median of Ratios


Gene	Sample A	Sample B
X	26	10
Y	26	10
Z	26	10
DE	2	50

Total = 80 80

Sample A Reads



Sample B Reads



Normalization: DESeq2 Median of Ratios

1. Take a row-wise average to produce an average sample (geometric mean) $\sqrt[n]{x_1 x_2 \cdots x_n}$

Gene	Sample A	Sample B		Avg. Sample
X	26	10	→	16
Y	26	10		16
Z	26	10		16
DE	2	50		10

Normalization: DESeq2 Median of Ratios

1. Take a row-wise average to produce an average sample (geometric mean) $\sqrt[n]{x_1 x_2 \cdots x_n}$

Gene	Sample A	Sample B	Avg. Sample
X	26	10	16
Y	26	10	16
Z	26	10	16
DE	2	50	10

2. Divide all rows by the Average Sample for that gene (**Ratio**)

Gene	Sample A/Avg.	Sample B /Avg.
X	26/16 = 1.6	10/16 = 0.6
Y	1.6	0.6
Z	1.6	0.6
DE	0.2	5

Normalization: DESeq2 Median of Ratios

1. Take a row-wise average to produce an average sample (geometric mean) $\sqrt[n]{x_1 x_2 \cdots x_n}$


Gene	Sample A	Sample B	Avg. Sample
X	26	10	16
Y	26	10	16
Z	26	10	16
DE	2	50	16

2. Divide all rows by the Average Sample for that gene (**Ratio**)

Gene	Sample A/Avg.	Sample B /Avg.
X	26/16 = 1.6	10/16 = 0.6
Y	1.6	0.6
Z	1.6	0.6
DE	0.2	5

3. Take the **median** of each column. Should be ~1 for all

Size factor	1.6	0.6
-------------	-----	-----



Normalization: DESeq2 Median of Ratios

1. Take a row-wise average to produce an average sample (geometric mean) $\sqrt[n]{x_1 x_2 \cdots x_n}$

Gene	Sample A	Sample B	Avg. Sample
X	26	10	16
Y	26	10	16
Z	26	10	16
DE	2	50	16

2. Divide all rows by the Average Sample for that gene (**Ratio**)
4. Divide all counts by sample specific size factor

Gene	Sample A/Avg.	Sample B /Avg.
X	26/16 = 1.6	10/16 = 0.6
Y	1.6	0.6
Z	1.6	0.6
DE	0.2	5

Gene	Sample A / S _A	Sample B / S _B
X	16.3	16.7
Y	16.3	16.7
Z	16.3	16.7
DE	1.3	83.3

Normalized counts for non-DE genes are similar!


3. Take the **median** of each column. Should be ~1 for all

Size factor	1.6	0.6
-------------	-----	-----

```
estimateSizeFactors(dds)
```


Assumption of DESeq2 Median of Ratios


Median of Ratios method assumes that most genes are not Differentially Expressed between samples.



Assumption of DESeq2 Median of Ratios

Median of Ratios method assumes that most genes are not Differentially Expressed between samples.

COUNTER EXAMPLE



NOTE: add back full picture or remove

- Late stage cell death (total RNA DOWN)
- High c-Myc cells (total RNA UP)


Known quantity spike-in transcripts (ERCC) can be used to normalize in these cases.

Normalization methods

Normalization method	Description	Accounted factors	Recommended use
CPM (counts per million)	$\frac{K_i}{Total\ Reads\ per\ Sample/10^6}$	sequencing depth	Comparison between replicates of the sample group
R/FPKM (reads/fragments per kilobase of exon per million reads/fragments mapped)	$\frac{K_i}{Gene\ Length/10^3 * Total\ Reads\ per\ Sample/10^6}$	sequencing depth and gene length	Comparison between genes in a sample
DESeq2's median of ratios [1]	K_i divided by sample-specific size factors	sequencing depth and RNA composition	Differential Expression between samples

Similar to DESeq2: EdgeR, limma-voom

Quality Control Visualizations



Examine sources of variation in the data

- Principal Component Analysis
- Hierarchical Clustering

(Log2 + 1) Transformed, Normalized Count Table

Gene	Sample A	Sample B	Sample C
1	1	1.6	0.5
2	2.2	-0.2	1
3	-1	1	3.1


Principle Component Analysis

Dimension reduction technique

Example: 3 gene dimensions -> 2 PC


Gene	Mock_12h	Mock_12h	Mock_24h	Mock_24h	HIV_12h	HIV_12h	HIV_24h	HIV_24h
Gene 1	8.9	8.9	8.9	9.0	8.9	8.9	9.0	6.8
Gene 2	0.6	-1.0	0.6	-1.0	0.6	-1.0	0.6	3.8
Gene 3	4.1	11.9	4.1	-0.5	4.1	8.7	4.0	4.4

original data space



PCA


component space



Do your samples cluster as expected?


What are the major sources of variation in the data?

Principle Component Analysis




- ✓ Do your samples cluster as expected?
- ✓ What are the major sources of variation in the data?


Principle Component Analysis




- ✓ Do your samples cluster as expected?
- ✓ What are the major sources of variation in the data?
- ✓ Is there a batch effect?



Differential Expression with DESeq2




Multi-factor experiment design



Factor 1:
Infection status (Mock or HIV)

Factor 2:
Time (12 or 24 hr)


Multi-factor experiment design



- Differential Expression compares two conditions
- We'll choose Infection status at 12 hr (Mock or HIV) for comparison
- We could also choose time, or a combination of multiple factors


Step 1: Modeling gene expression values

All leading DE tools use **regression models** to estimate the fold change between conditions for **each gene**



Step 1: Modeling gene expression values

All leading DE tools use **regression models** to estimate the fold change between conditions for **each gene**
Example, simple linear regression:



$$Y = \beta_0 + \beta_1 X + e$$

Log2 Expression Values

Intercept

Condition (0-Mock, 1-HIV)


Slope: difference between Mock / HIV

Error

DESeq2 uses a Generalized Linear Model with a Negative Binomial error Distribution, which has been shown to be best fit for RNAseq data.

Love, M.I., Huber, W. & Anders, S. Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2. *Genome Biol* **15**, 550 (2014).

Step 2: Hypothesis Testing



$$H_0 : \beta_1 = 0 \quad \text{vs.} \quad H_A : \beta_1 \neq 0$$

H_0 : there is no systematic difference between the average read count values for Mock vs. HIV

- Statistical test – Wald test (similar to t-test) on β_1
- $Z = \beta_1 / SE_{\beta_1}$
- Z-statistic is compared to the normal distribution and probability of getting a statistic at least as extreme is computed

Is EGR1 differentially expressed?

Yes! $p \ll 0.05$

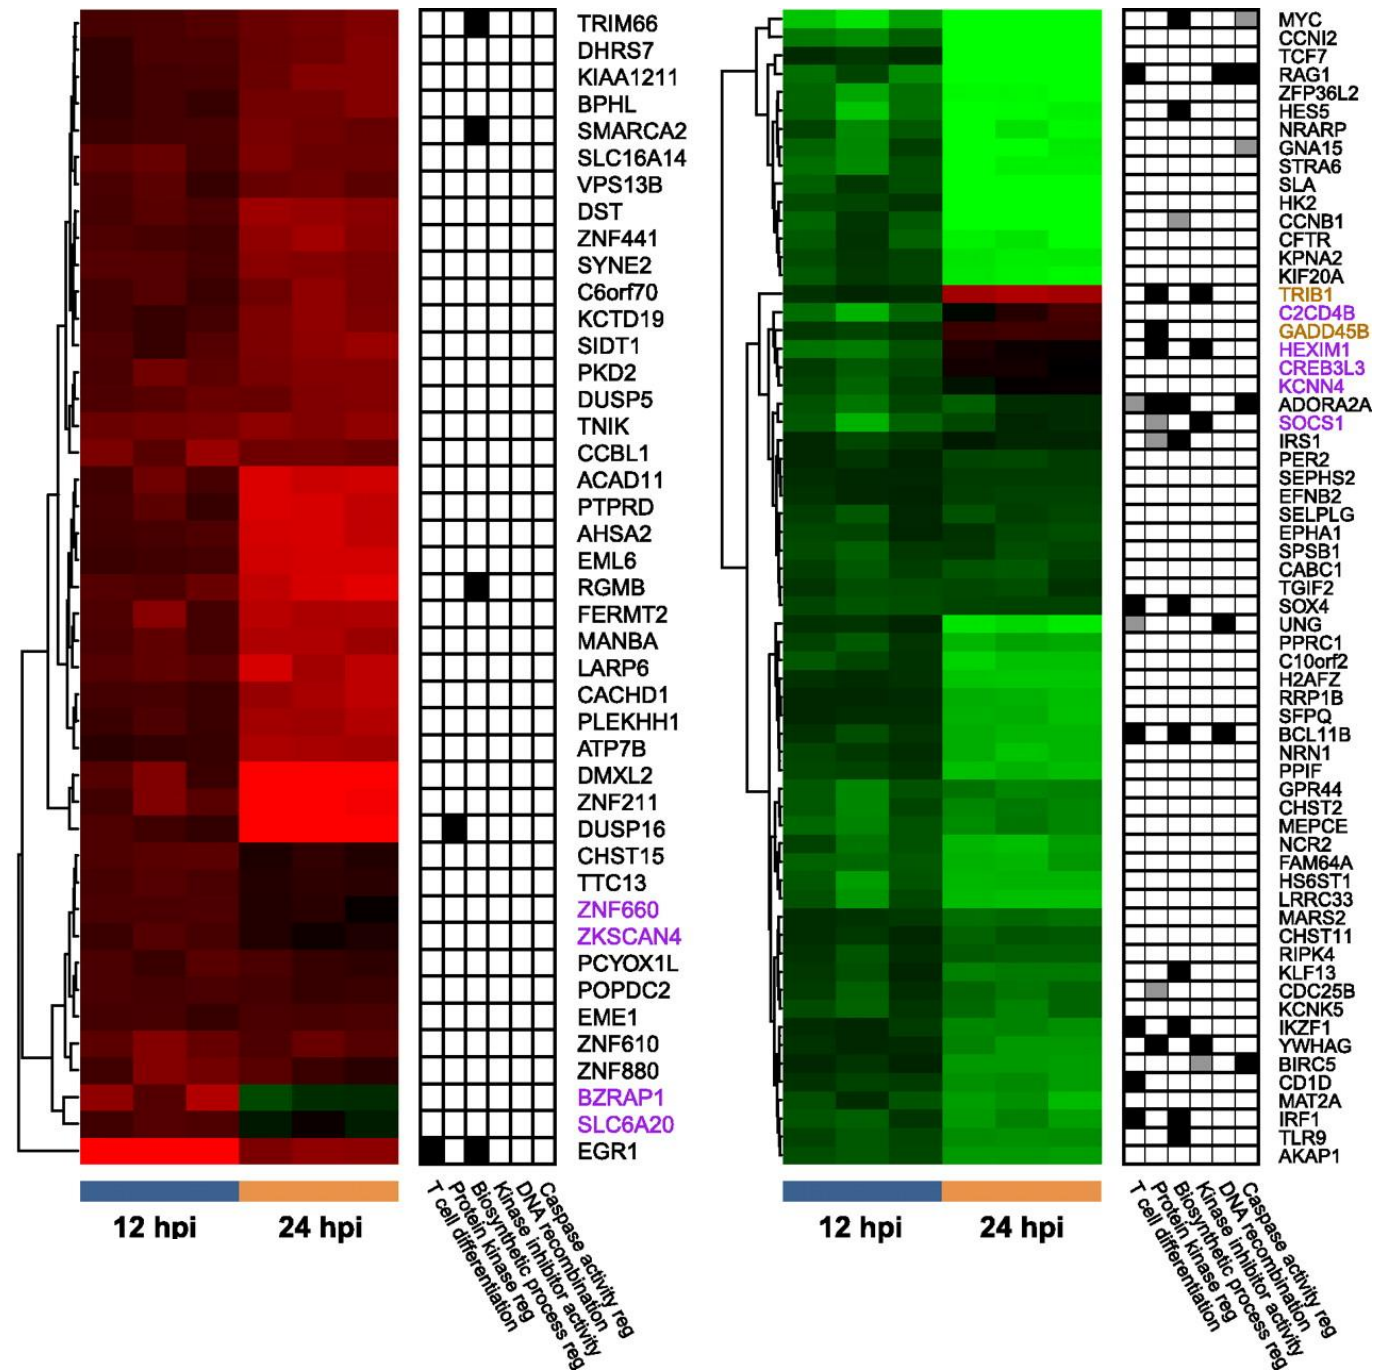

DESeq2 Results table

GeneID	Base mean	log2FoldChange	StdErr	P-value	P-adj
EGR1	1273	1.55	0.13	1.19e-77	1.52e-73
MYC	5226	-1.53	0.14	1.63e-36	1.03e-32

- Mean of normalized counts – averaged over all samples from two conditions (HIV, Mock)
- Log of the fold change between two conditions
- StdErr – Standard error of coefficient (e.g. b_1)
- P-value – the probability that the Wald statistic is as extreme as observed if H_0 were true
- P-adj – accounting for multiple testing correction

Study findings

- T cell differentiation-related genes were overrepresented in the DEG at 24hr
- 'Large-scale disruptions to host transcription' at 24hr



References

DESeq2 vignette (R/Rstudio):

<http://www.bioconductor.org/packages/release/bioc/vignettes/DESeq2/inst/doc/DESeq2.html#differential-expression-analysis>

HBC Training (Command line/R):

https://hbctraining.github.io/DGE_workshop

Galaxy Training:

https://galaxyproject.org/tutorials/rb_rnaseq/